

Oslo Bioinformatics Workshop Week 2025

Statistical principles in machine learning for small biomedical data

Manuela Zucknick

Oslo Centre for Biostatistics and Epidemiology, University of Oslo
manuela.zucknick@medisin.uio.no

December 12, 2025

Some of the figures in this presentation are taken from "An Introduction to Statistical Learning, with applications in R" (Springer, 2021) with permission from the authors.

Schedule for Today

Schedule

Time	Topic	Presenter
Now	<u>Preparations</u>	
09:00 - 10:00	<u>(Supervised) machine learning with small data</u>	Manuela Zucknick
	<u>R lab 1</u>	Manuela Zucknick
10:15 - 11:15	<u>Overfitting, regularisation and all that</u>	Manuela Zucknick
	<u>R lab 2</u>	Manuela Zucknick
11:30 - 12:00	<u>The potential of Bayesian modelling for complex biomedical data</u>	Manuela Zucknick

Github, Workshop webpage and Posit Cloud project

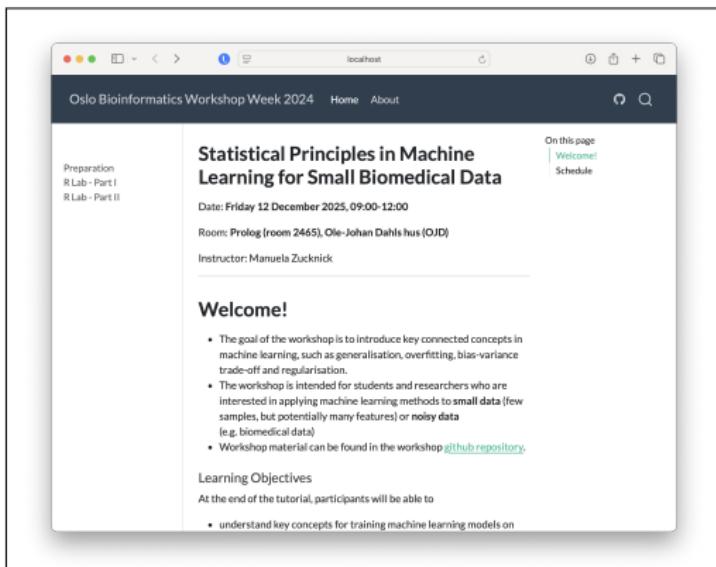
- Github:

<https://github.com/ocbe-uio/workshop-stat-higdim/>

- Workshop webpage:

<https://ocbe-uio.github.io/workshop-stat-higdim/>

- Posit Cloud project: <https://posit.cloud/content/5131383/>



- PhD course at MedFak, UiO (Valeria Vitelli, OCBE)

The screenshot shows a web browser window for the University of Oslo (ui.no) with the following details:

- Page Title:** UNIVERSITY OF OSLO
- Language:** NO EN
- Section:** Studies > Courses
- Course Information:** MF9385 – Introduction to machine learning in biomedical research
- Credit:** Credits: 5
- Level:** Level: PhD
- Teaching:** Spring Spring 2026: StudentWeb opens for registration 1.12.2025. Teaching dates and sign-up deadline is posted on the semester page.
- Examination:** Spring Teaching language: English

Course description

→ Course content → Overlapping courses
→ Learning outcome → Teaching
→ Admission to the course → Examination

- OCBE Statistical Advising Service
- <https://www.med.uio.no/imb/english/research/centres/ocbe/advising/>

Some topics for this morning

Part 1

- What is supervised machine learning?
- What do we mean by small data?
- What can we do to improve ML with small data?
 - Restrict the model space → Regularisation
 - Borrow information → Include known structure in the model

Part 2

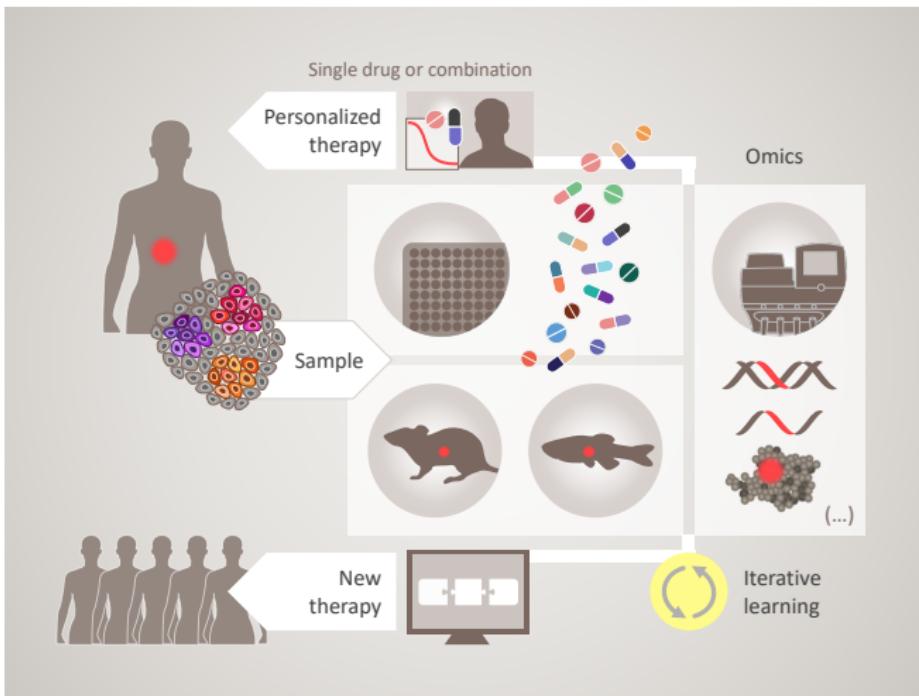
- Overfitting
- Variance vs bias
- Model selection, assessment & validation
- Prediction performance
- Resampling: Cross-validation

Introductory example:

Integrative omics for personalized cancer therapy

Personalized cancer therapy

...aims to find the best therapy for each patient based on data about the patient and tumor (e.g. genomic data).



Predict sensitivity to multiple drugs \mathbf{Y} from multi-omics \mathbf{X}

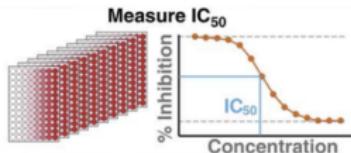
$$\mathbf{Y} = \mathbf{XB} + \epsilon$$

- **Multivariate \mathbf{Y} :**

Drug dose response

drug sensitivity

$$n \text{ cell lines} \left[\begin{array}{c|c|c} & \dots & \\ \text{y}_{\bullet 1} & \dots & \text{y}_{\bullet m} \\ & \dots & \end{array} \right] = \mathbf{Y}$$

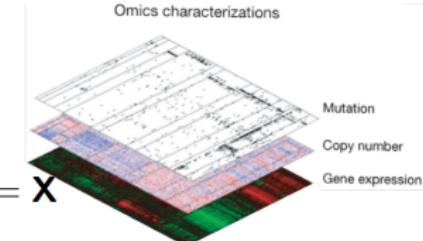


Source: Yang, et al. 2017

- **Heterogeneous \mathbf{X} :**

Integrative omics

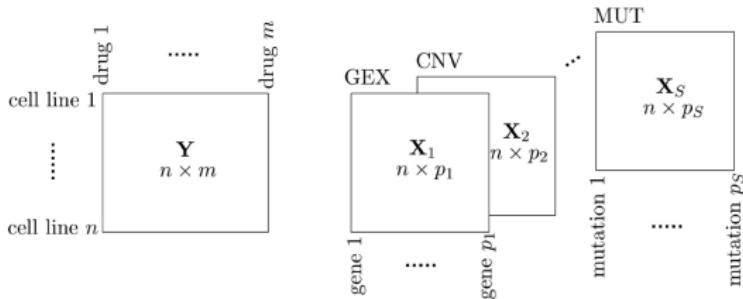
$$n \text{ cell lines} \left[\begin{array}{c|c|c} \text{gene expression} & \text{copy number} & \text{mutation} \\ \hline \text{X}_1 & \text{X}_2 & \text{X}_3 \\ | & | & | \\ \end{array} \right] = \mathbf{X}$$



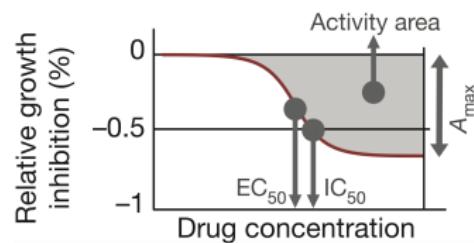
Source: TCGA, 2013

Challenges and opportunities (1)

- Small sample size
- Several types of input data \mathbf{X} :
E.g., gene expression, copy number, mutation
- Multivariate response \mathbf{Y}



- Unclear how to define \mathbf{Y}



Challenges and opportunities (2)

The data are highly **structured**:

1. **In Y:** relationships between drugs, e.g. due to similar chemical drug composition, same target genes/pathways
2. **In X:** relationships between molecular data sources

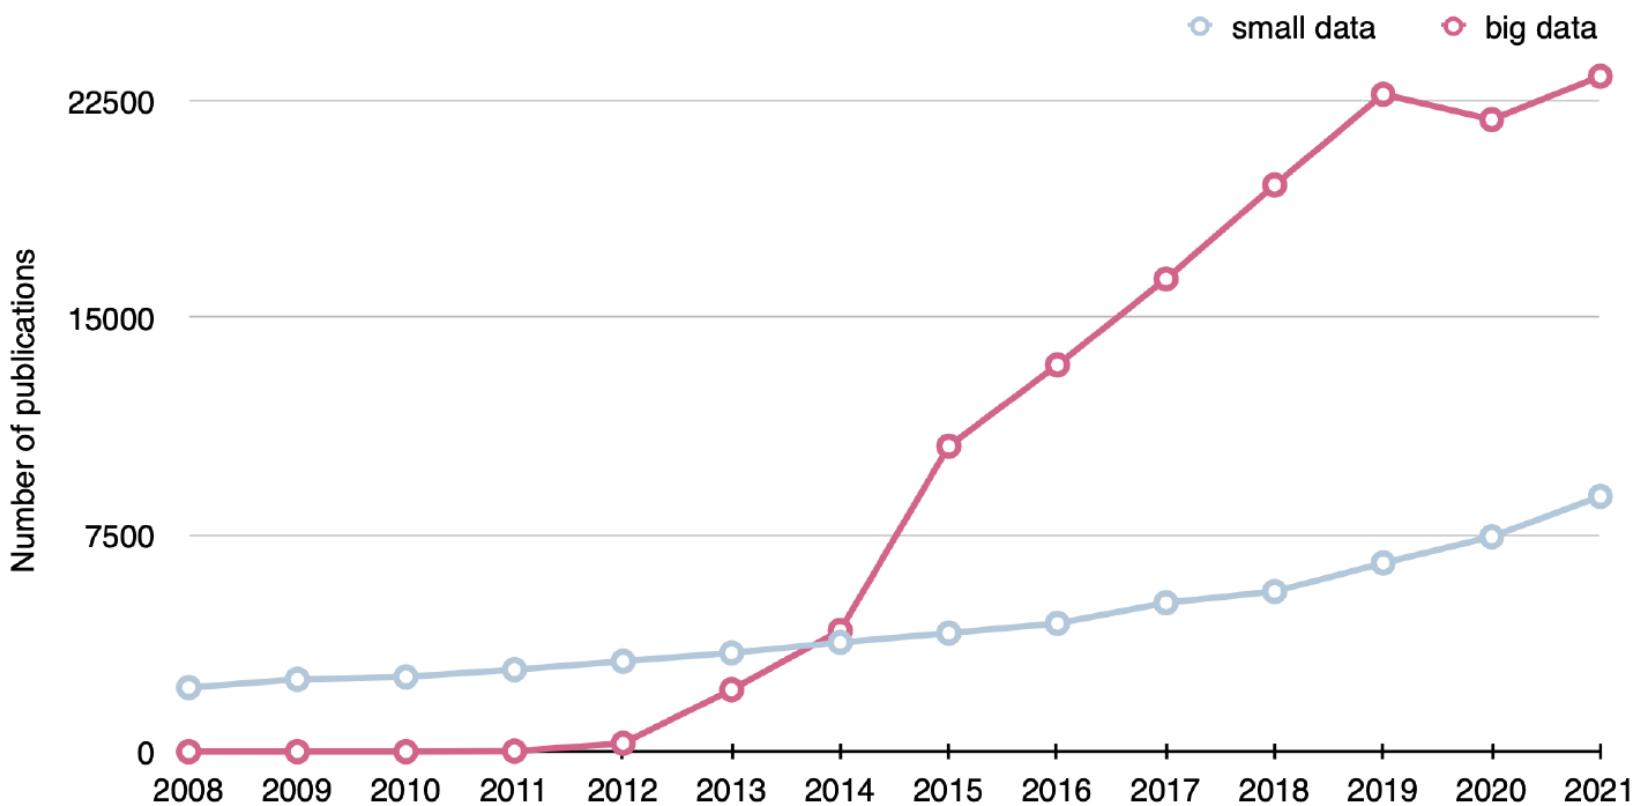
a	Function	Memory	Environment	Message	Product	Result
b	Central dogma of molecular biology	Genome (DNA)	Epigenome and other regulatory elements (e.g. chromatin modifications, miRNA, TFs) 	Transcriptome (mRNA) 	Proteome (protein) 	Phenome (cell, tissue, organism) 
c	Data types	CN, SNPs, LOH 	Histone modification TF binding, miRNA, methylation 	GE 	Protein expression 	Phenotype, clinical characteristics 

Ickstadt et al. (2018)

(Supervised) Machine Learning with Small Data

Manuela Zucknick (with slides from Maren Hackenberg)

Machine learning with small data



Machine learning with small data

- What do we mean by “small data”?
 - Implications for machine learning?
 - Aspects when building (multi-omic) machine learning predictors of drug response (e.g. Sammut et al. Nature 2022):
 1. Biological knowledge +
 2. Feature selection +
 3. Prioritisation of accessible data types +
 4. Machine learning algorithms
- Develop ML methods that allow us to consider aspects 1 to 3.

What is supervised machine learning?

Supervised learning

refers to the task of inferring a functional relationship between **input data matrix \mathbf{X}** (e.g. gene expression array measurements) and **output data vector Y** (= response/ outcome).

The input data are used for **predicting** the outcome.

$$Y = f_{\beta}(\mathbf{X}) + \epsilon,$$

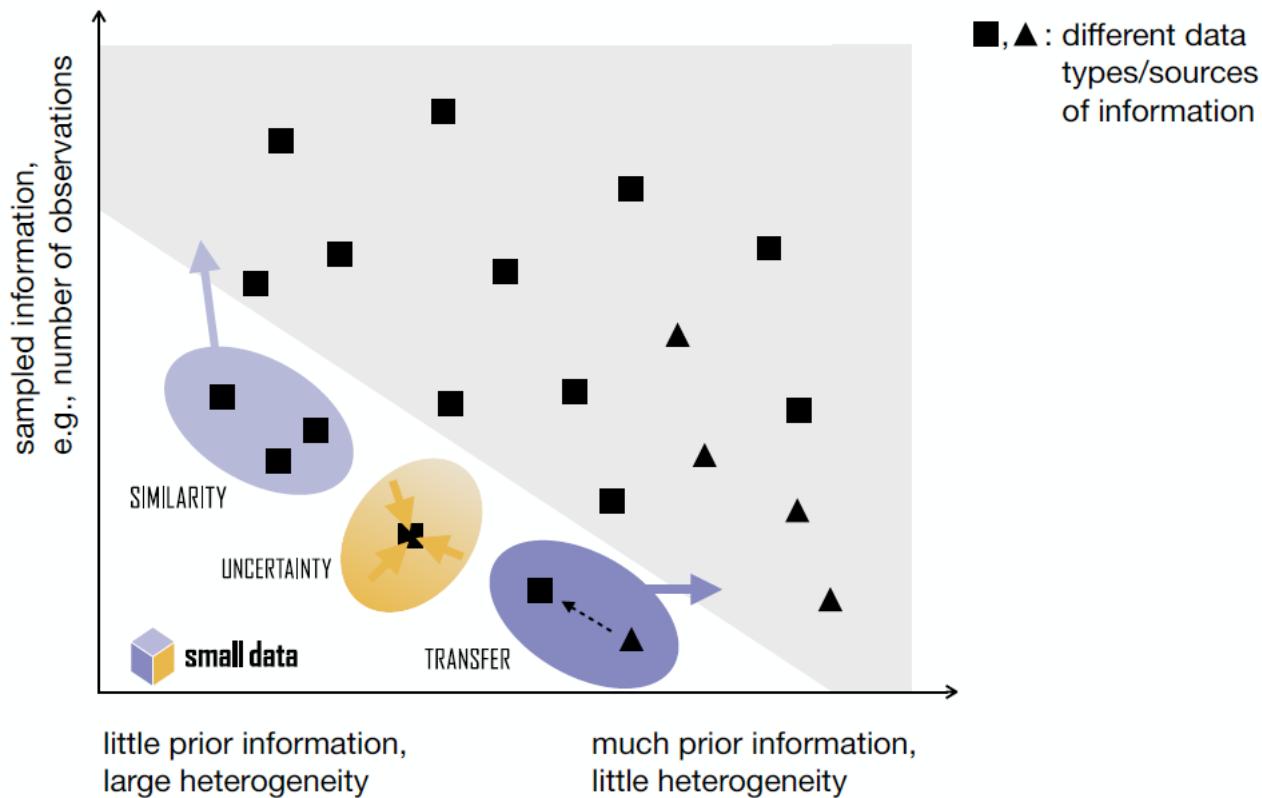
where ϵ captures measurement errors and other discrepancies, e.g. by $\epsilon \sim N(0, \sigma^2 I_n)$.

In classical statistics, this task is usually performed by **(generalised) linear regression models**.

What do we mean by small data?

- Large p, small n ($p > n$)
- Potentially, more variables in the model than we have samples
- Classical statistical methods (e.g. linear regression) do not work:
- More parameters (e.g. regression coefficients) to estimate than observations for estimating them
- Even if all parameters can be estimated: Danger of over-fitting
- Example: Predict treatment response using gene expression data
($n \sim 100$, $p \sim 20000$)

What do we mean by small data?



What can we do?

- (1) Restrict the model space
- (2) Borrow information across observations
- (3) Increase sample size ☺

What can we do?

(1) Restrict the model space

- (A) Careful feature engineering:
 - Preselect variables by biological relevance
 - Non-specific filtering, e.g. keep only variables with variance across observations larger than a threshold
- (B) Make use of known structure in the data (biological knowledge)
- (C) Use of regularisation techniques:
 - L1 and L2 penalisation
 - add a penalty term to the loss function to reduce the complexity of the model
 - Bayesian equivalents: restrictions on the prior (Bayesian variable selection)
 - Early stopping
 - train a model iteratively only until the validation error starts to decrease (boosting, neural networks)
 - Dropout regularisation
 - randomly dropping out neurons while training (neural networks) or
 - randomly dropping features when building a regression tree (random forest)

Penalised regression

- Standard regression cannot deal with $p \gg n$:
 - The maximum-likelihood estimate $\hat{\beta} = \arg \max_{\beta} \ell(\beta)$ does not exist ($\ell = \log\text{-likelihood}$).
- **Solution:**
Penalise the likelihood function by subtracting a penalty term and maximise penalised log-likelihood instead:

$$\hat{\beta} = \arg \max_{\beta} (\ell(\beta) - \lambda \|\beta\|)$$

- λ is a **penalty parameter**,
- $\|\beta\|$ represents the size of the regression coefficient vector,
- The larger λ is chosen, the more the algorithm is encouraged to find a solution where $\|\beta\|$ is small \rightarrow **shrinkage**.

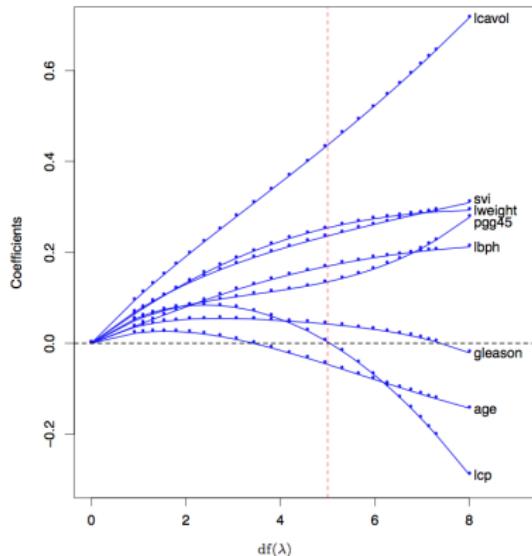
Penalised regression

- Examples for penalty terms:
 - Ridge regression (Hoerl and Kennard 1970):
$$\lambda \|\beta\| := \lambda \sum_{g=1}^p \beta_g^2 \quad \rightarrow \mathbf{L}_2 \text{ penalty}$$
 - Lasso regression (Tibshirani 1996):
$$\lambda \|\beta\| := \lambda \sum_{g=1}^p |\beta_g| \quad \rightarrow \mathbf{L}_1 \text{ penalty}$$
 - Elastic net (Zou and Hastie 2005):
Combination of both ridge and lasso penalty:
$$\lambda_1 \sum_{g=1}^p |\beta_g| + \lambda_2 \sum_{g=1}^p \beta_g^2$$
- Advantage of lasso and elastic net:
Both will produce a sparse solution, where only a few genes have estimate $\hat{\beta}_g \neq 0$.

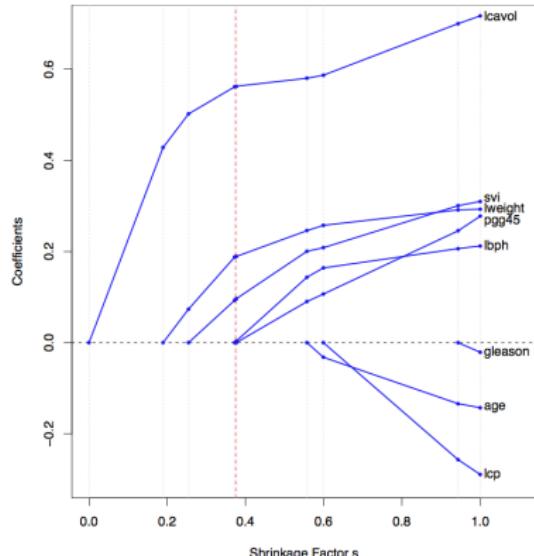
Penalised regression

Examples for coefficient paths relative to penalty λ :

Ridge regression



Lasso regression



Hastie et al. (2009), Figures 3.8 and 3.10

Penalised regression

- Ridge regression L_2 : shrinks all coefficients to small, but non-zero values.
- Lasso regression L_1 : shrinks some coefficients to exactly zero.
- Elastic net: mixture of the two: does shrink some coefficients to exactly zero. Keeps more variables if there is correlation.

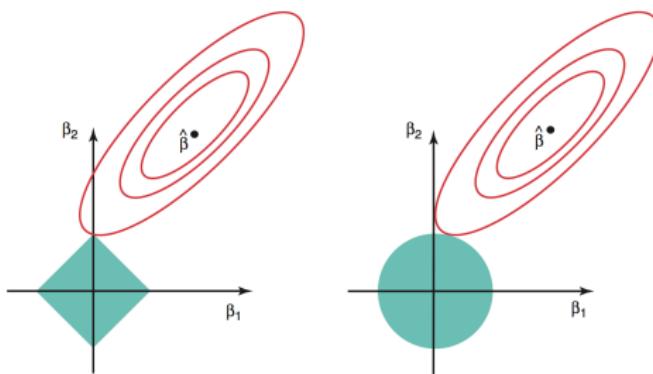


FIGURE 6.7. Contours of the error and constraint functions for the lasso (left) and ridge regression (right). The solid blue areas are the constraint regions, $|\beta_1| + |\beta_2| \leq s$ and $\beta_1^2 + \beta_2^2 \leq s$, while the red ellipses are the contours of the RSS.

Different penalties for different types of data

Assume two data matrices \mathbf{X} and \mathbf{Z} :

$$\mathbf{Y} = \mathbf{X}\beta + \mathbf{Z}\gamma + \epsilon$$

- **Mandatory covariates:** Do not penalise the parameters γ :

$$\ell_{\text{pen}}(\beta, \gamma) = \ell(\beta, \gamma) - \lambda \|\beta\|$$

e.g. with R packages `glmnet` or `penalized`

- **Several types of molecular data sets:**

Allow different penalties for β and γ :

$$\ell_{\text{pen}}(\beta, \gamma) = \ell(\beta, \gamma) - \lambda_\beta \|\beta\| - \lambda_\gamma \|\gamma\|$$

e.g. with R packages `GRridge` (Van de Wiel *et al.*, 2016)
<http://www.few.vu.nl/~mavdwiel/grridge.html>)

Different penalties for different types of data

Assume two data matrices \mathbf{X} and \mathbf{Z} :

$$Y = \mathbf{X}\beta + \mathbf{Z}\gamma + \epsilon$$

- Several types of molecular data sets:
- Alternative: **Combine all data and use one penalty**, after scaling all features to unit variance to ensure that the data sources are treated equally.
- Example: Elastic Net models in Barretina et al. (2012)

Bayesian interpretation of penalised regression

- Ridge regression as a penalised log-likelihood problem ...

$$\hat{\beta} = \arg \max_{\beta} (\ell - \lambda \sum_{i=1}^p \beta_i^2)$$

- ... is equivalent to *maximum a posteriori* solution of Bayesian linear regression with Gaussian prior

$$p(\beta|\tau) = N(0, \tau I_p), \text{ where } \tau = 1/(2\lambda) :$$

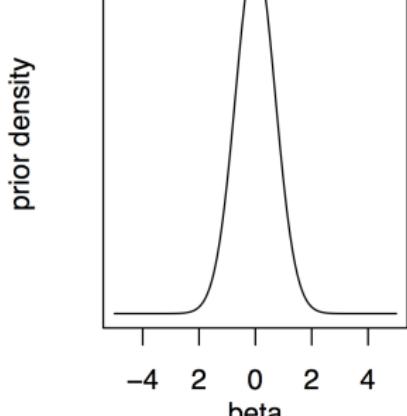
$$\begin{aligned} \text{posterior} &\propto \text{likelihood} \times \text{prior} \\ p(\beta|X, Y, \tau) &\propto p(Y|\beta, X)p(\beta|\tau) \\ \Leftrightarrow \log(\beta|X, Y, \tau) &\propto \ell + \log p(\beta|\tau) \\ &\propto \ell - \lambda \sum_{i=1}^p \beta_i^2 - C \end{aligned}$$

Bayesian interpretation of penalised regression

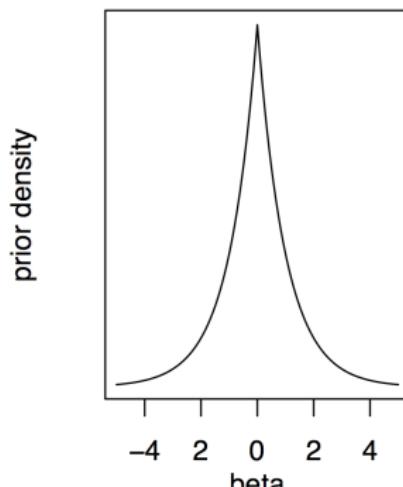
- Generalisation: Bridge regression (e. g. Frank and Friedman 1993)

$$\hat{\beta} = \arg \max_{\beta} (\ell - \lambda \sum_{i=1}^p |\beta_i|^q) \quad (q > 0)$$

Ridge ($q = 2$): Gaussian prior



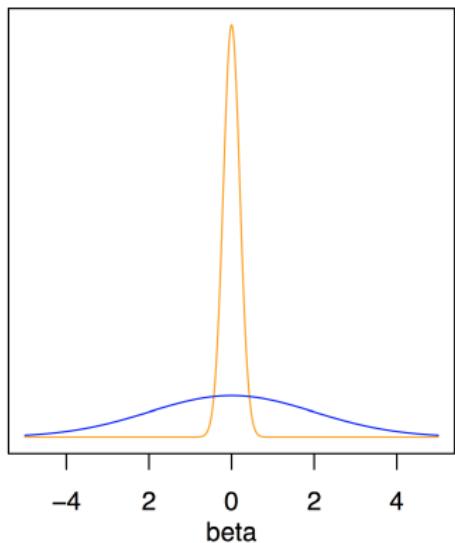
Lasso ($q = 1$): Laplace prior



Bayesian hierarchical model for variable selection (BVS)

- Bayesian variable selection model with indicator variable

$$\gamma_i = \begin{cases} 1 & \text{variable } i \text{ is included} \\ 0 & \text{variable } i \text{ is excluded} \end{cases}$$



E. g. normal mixture prior
(George and McCulloch 1993):

$$\beta_i | \gamma_i \sim (1-\gamma_i)N(0, \sigma^2) + \gamma_i N(0, g\sigma^2)$$

where $\sigma^2 > 0, g > 0$

Bayesian hierarchical model for variable selection (BVS)

Advantages:

- Flexibility in penalization through large variety of possible prior distributions for β (hierarchical models)
- Full posterior distributions, including probabilities for the selection of variables and posterior distributions for β ;
- Improve prediction performance by Bayesian Model Averaging

Bayesian Model Averaging for Linear Regression Models

Adrian E. Raftery; David Madigan; Jennifer A. Hoeting

Journal of the American Statistical Association, Vol. 92, No. 437 (Mar., 1997), 179-191.

Stable URL:
<http://links.jstor.org/sici?&sici=0162-1459%28199703%2992%3A437%3C179%3ABMAFLR%3E2.0.CO%3B2-9>

Journal of the American Statistical Association is currently published by American Statistical Association.

The JSTOR logo consists of a stylized letter 'J' enclosed within a square frame, with the word 'STOR' in a serif font below it and a registered trademark symbol (®) to the right.

What can we do?

(2) Borrow information

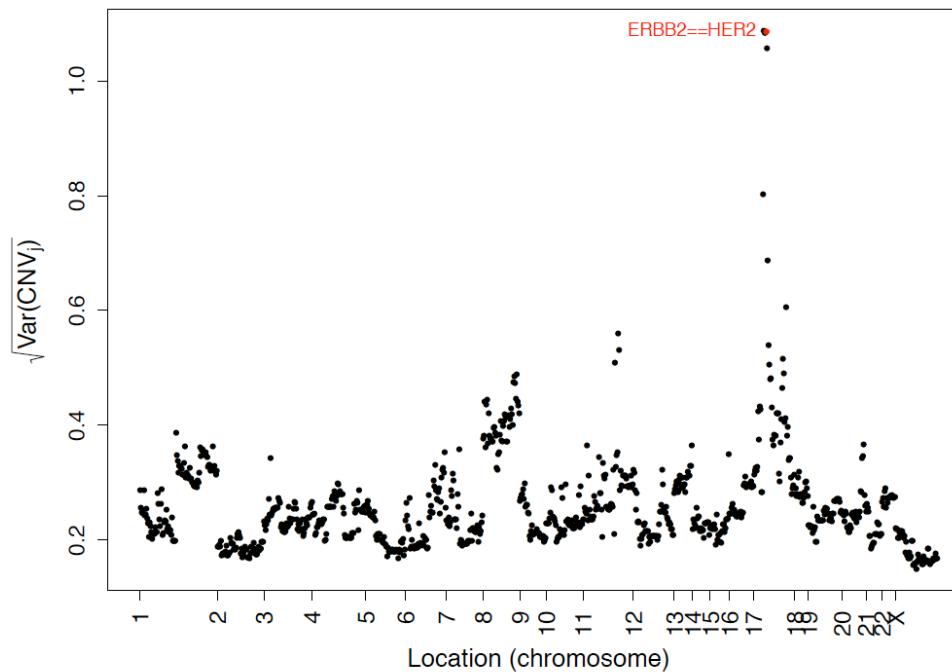
- Borrow information across observations in the data set
- If there is correlation, include this in your model
 - between variables (e.g. MRF prior for defining which variables to include together)
 - between samples (covariance matrix)
- Borrow information from external knowledge
 - E.g., use pathways to determine which genes should be included together
- Borrow information across data sets: transfer learning

Make use of external (biological) knowledge

- (1) Use known relationships with one data source (CNV) to guide the variable selection in another (gene expression)
- (2) Combine the data-driven ML approach with knowledge-driven mechanistic modelling
- (3) Make use of correlations in the data
 - between input variables - to restrict the model space
 - between response variables - to borrow information

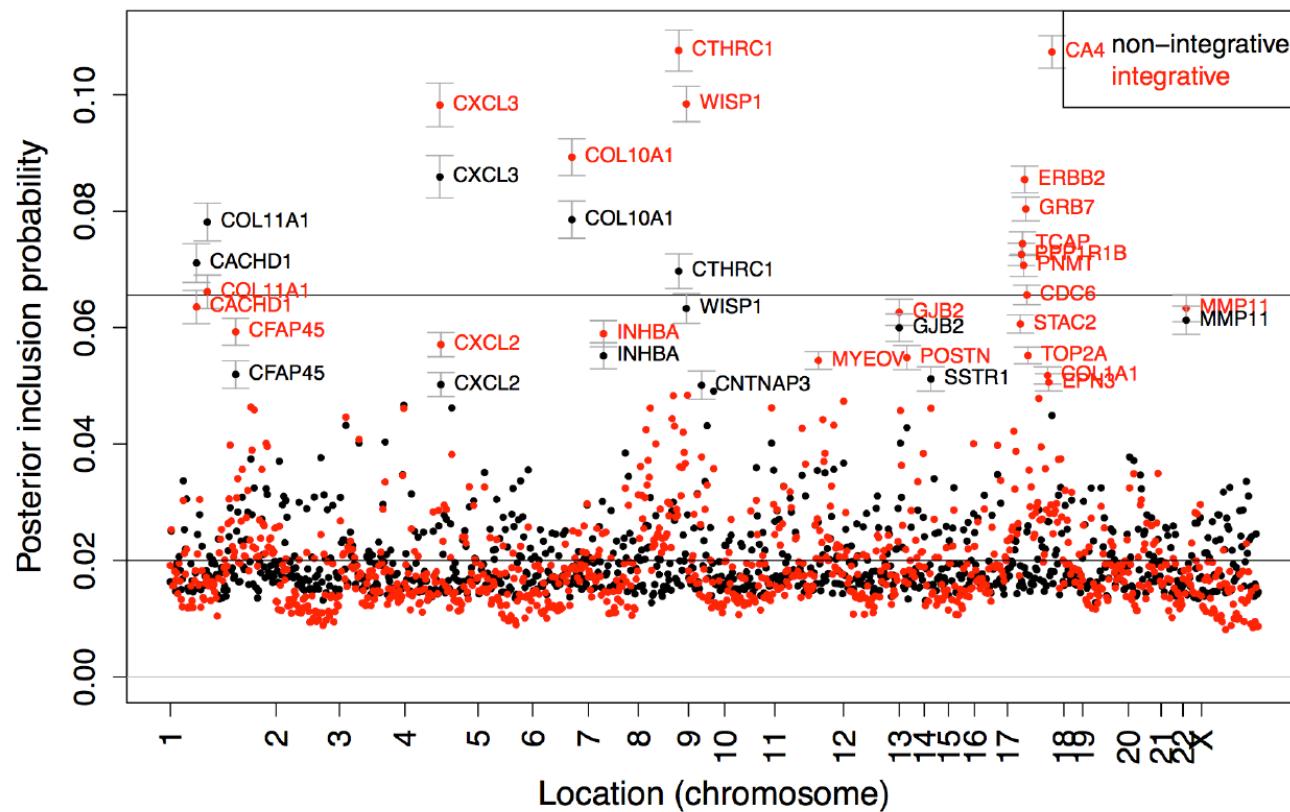
(1) Use known relationships with one data source to guide the variable selection in another

Std. dev. of CNV data of HER2-pos. breast cancer and healthy tissue samples



Idea: Use CNV information to weigh prior inclusion probabilities of gene expression variables in Bayesian variable selection

(1) Use known relationships with one data source to guide the variable selection in another



HER2 (= ERBB2) only selected in integrative analysis

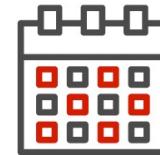
(2) Combine the data-driven ML approach with knowledge-driven mechanistic modelling

An exemplary small data challenge: Learn disease trajectories of patients with spinal muscular atrophy



Baseline characterisation

- age
- SMA subtype
- ...



Different motor function tests over time

- RULM
- HFMSE
- ...



Latent health status

$$\frac{d}{dt} \mu(t) = ?$$

Explicit model



Subgroup-specific local models



Heterogeneity



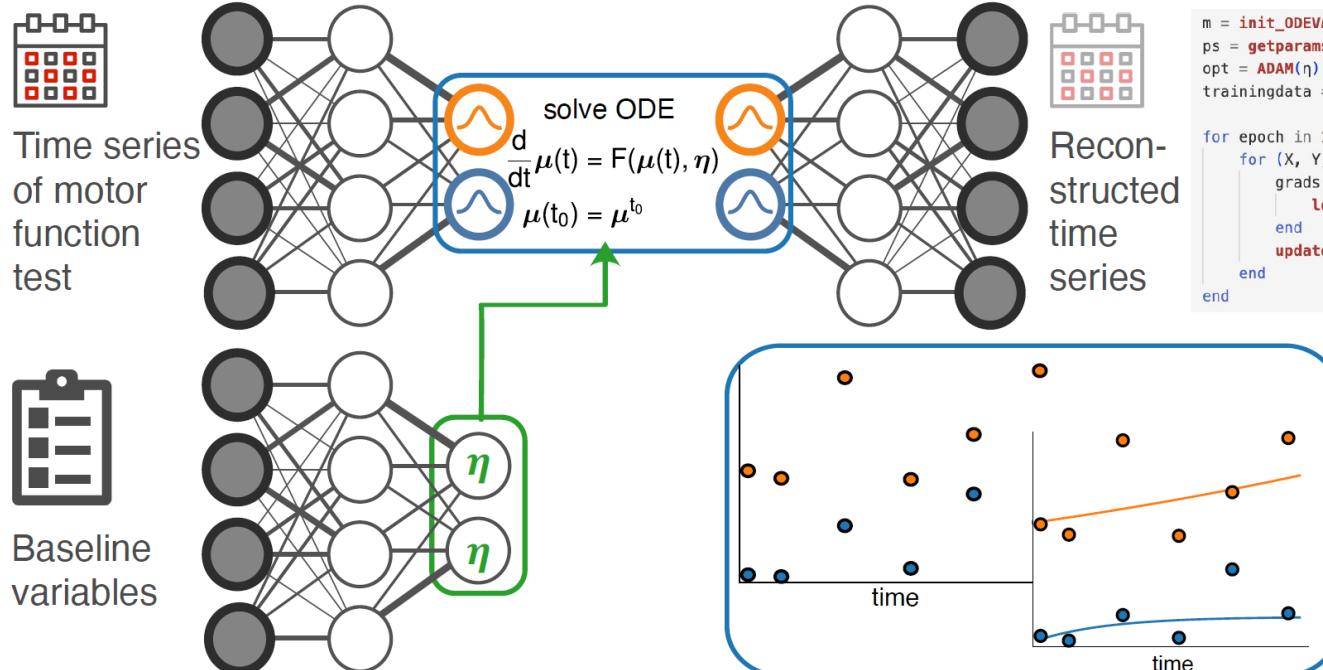
Irregular time points



RULM
HFMSE
Different motor function tests

(2) Combine the data-driven ML approach with knowledge-driven mechanistic modelling

Describe individual SMA trajectories as ODEs in the latent space of a deep learning model



```
m = init_ODEVAE()
ps = getparams(m)
opt = ADAM(η)
trainingdata = zip(xs, xs_baseline, tvals)

for epoch in 1:epochs
    for (X, Y, t) in trainingdata
        grads = gradient(ps) do
            loss(X, Y, t, m, args=args)
        end
        update!(opt, ps, grads)
    end
end
```

(3) Make use of correlations in the data: between input variables - to restrict the model space

BayesSUR: An R Package for High-Dimensional Multivariate Bayesian Variable and Covariance Selection in Linear Regression

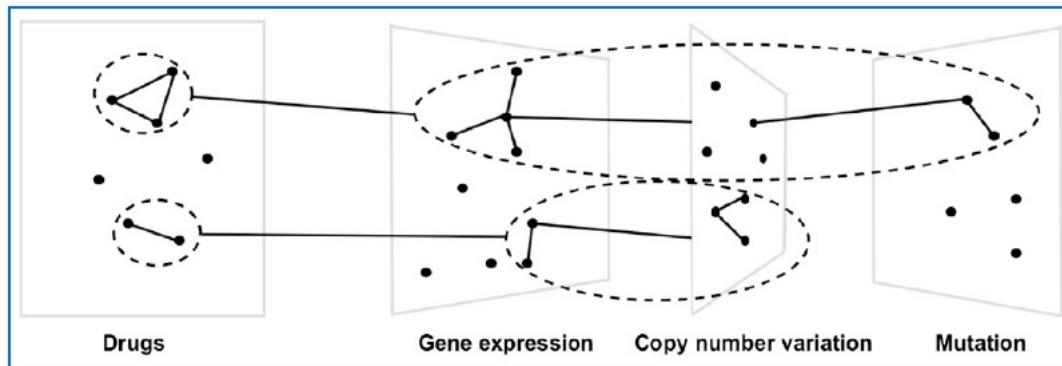
Zhi Zhao, Marco Banterle, Leonardo Bottolo, Sylvia Richardson, Alex Lewin, Manuela Zucknick

Vol. 100, Issue 11

 Paper

 R package (BayesSUR)

 R replication code



-> See later

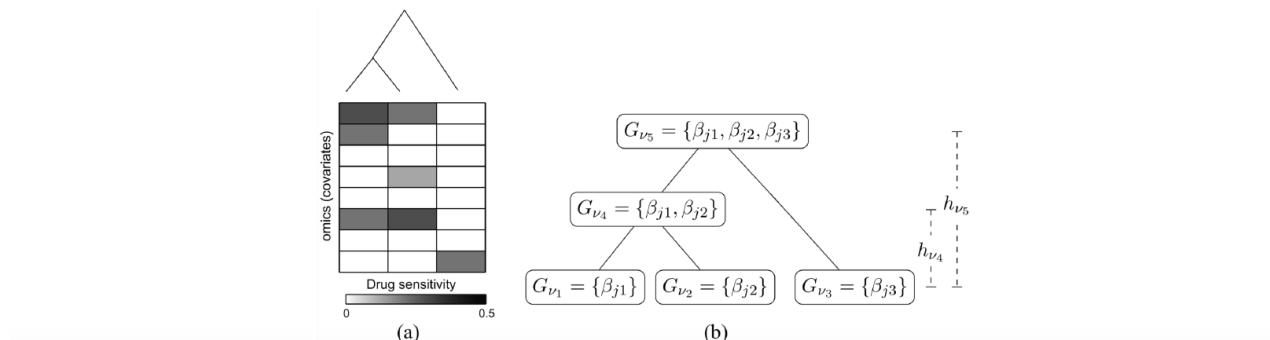
(3) Make use of correlations in the data: between response variables - to borrow information

(Multi-response) Tree-guided group lasso (Kim & Xing 2012)

- Include dependencies between columns of \mathbf{Y} in a group lasso
- Extension to IPF-tree lasso

$$\text{Tree lasso: } \text{pen}(\mathbf{B}) = \lambda \sum_{j=1}^p \sum_{\nu \in \{V_{\text{int}}, V_{\text{leaf}}\}} \omega_\nu \|\beta_j^{G_\nu}\|_{\ell_2}$$

$$\text{IPF-tree lasso: } \text{pen}(\mathbf{B}) = \sum_s \lambda_s \left(\sum_{j_s} \sum_{\nu \in \{V_{\text{int}}, V_{\text{leaf}}\}} \omega_\nu \|\beta_{j_s}^{G_\nu}\|_{\ell_2} \right)$$



(3) Make use of correlations in the data: between response variables - to borrow information

Drug screens for precision cancer medicine: How to predict the drugs' effect with data on drugs and tumour?

ROYAL STATISTICAL SOCIETY
DATA | EVIDENCE | DECISIONS

Journal of the Royal Statistical Society
Applied Statistics
Series C

Original Article | Open Access | CC BY SA

Structured penalized regression for drug sensitivity prediction

Zhi Zhao, Manuela Zucknick

